

ADS 505 Final Project

Team 1

Angela Zhang

Emina Beleanich

Luke Awino

A stroke occurs when blood supply is blocked or a blood vessel bursts in the brain. Strokes can cause brain damage, long-term disabilities and in the most severe cases, death. According to the Centers for Disease Control and Prevention, someone in the United States will suffer from a stroke every forty seconds. Along with cancer and heart disease, strokes are a leading cause of death for Americans. There are a wide variety of factors that can contribute to a person’s likelihood of suffering from a stroke. Some risk factors include smoking, obesity, age and high blood pressure (Centers for Disease Control and Prevention [CDC], 2021). Stroke prediction is essential in saving lives. Therefore, it is important that we understand the causes of a stroke in order to make better recommendations to those at risk through the utilization of predictive modeling.

```
In [1]: #import the packages
%matplotlib inline
%config IPCompleter.greedy=True

from pathlib import Path

import pandas as pd
import numpy as np
from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
import matplotlib.pyplot as plt
from sklearn import preprocessing
from dmbs import plotDecisionTree, classificationSummary, regressionSummary
from sklearn.linear_model import LinearRegression, Lasso, Ridge, LassoCV, BayesianRidge, LogisticRegressionCV, LogisticRegression
import statsmodels.formula.api as sm
import matplotlib.pyplot as plt
import matplotlib as mpl
import matplotlib.pyplot as plt
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
import seaborn as sns
from dmbs import gainsChart
from dmbs import regressionSummary, exhaustive_search
from dmbs import backward_elimination, forward_selection, stepwise_selection
from dmbs import adjusted_r2_score, AIC_score, BIC_score
from sklearn.metrics import classification_report
from sklearn.naive_bayes import MultinomialNB
```

```
In [2]: #Load the dataset
stroke_df = pd.read_csv('/Users/Luke-Workstation/Desktop/ADS 505/datasets/healthcare-dataset-stroke-data.csv')
stroke_df
```

Out[2]:

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	stroke
0	9046	Male	67.0	0	1	Yes	Private	Urban	228.69	36.6	formerly smoked	0
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	202.21	NaN	never smoked	0
2	31112	Male	80.0	0	1	Yes	Private	Rural	105.92	32.5	never smoked	0
3	60182	Female	49.0	0	0	Yes	Private	Urban	171.23	34.4	smokes	0
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural	174.12	24.0	never smoked	0
...
5105	18234	Female	80.0	1	0	Yes	Private	Urban	83.75	NaN	never smoked	0
5106	44873	Female	81.0	0	0	Yes	Self-employed	Urban	125.20	40.0	never smoked	0
5107	19723	Female	35.0	0	0	Yes	Self-employed	Rural	82.99	30.6	never smoked	0
5108	37544	Male	51.0	0	0	Yes	Private	Rural	166.29	25.6	formerly smoked	0
5109	44679	Female	44.0	0	0	Yes	Govt_job	Urban	85.28	26.2	Unknown	0

5110 rows × 12 columns

Inspecting the raw data is an important step in optimizing the best models. We begin our exploratory data analysis by importing relevant packages and the chosen dataset into the environment. Following this step, the structure of the dataset is acquired.

```
In [3]: stroke_df.shape
```

Out[3]: (5110, 12)

```
In [4]: stroke_df.columns
```

Out[4]: Index(['id', 'gender', 'age', 'hypertension', 'heart_disease', 'ever_married', 'work_type', 'Residence_type', 'avg_glucose_level', 'bmi', 'smoking_status', 'stroke'], dtype='object')

Inspecting statistical measures for numerical variables allows us to detect outliers and handle missing data.

```
In [5]: stroke_df.describe() #summary statistics for the numerical variables
```

Out[5]:

	id	age	hypertension	heart_disease	avg_glucose_level	bmi	stroke
count	5110.000000	5110.000000	5110.000000	5110.000000	5110.000000	4909.000000	5110.000000
mean	36517.829354	43.226614	0.097456	0.054012	106.147677	28.893237	0.048728
std	21161.721625	22.612647	0.296607	0.226063	45.283560	7.854067	0.215320
min	67.000000	0.080000	0.000000	0.000000	55.120000	10.300000	0.000000
25%	17741.250000	25.000000	0.000000	0.000000	77.245000	23.500000	0.000000
50%	36932.000000	45.000000	0.000000	0.000000	91.885000	28.100000	0.000000
75%	54682.000000	61.000000	0.000000	0.000000	114.090000	33.100000	0.000000
max	72940.000000	82.000000	1.000000	1.000000	271.740000	97.600000	1.000000

```
In [6]: stroke_df.dtypes # check the data types
```

Out[6]: id int64
gender object
age float64
hypertension int64
heart_disease int64
ever_married object
work_type object
Residence_type object
avg_glucose_level float64
bmi float64
smoking_status object
stroke int64
dtype: object

```
In [7]: counts = stroke_df.nunique() #check for unique counts. need to check gender for the three different counts, work type counts
```

Out[7]: id 5110
gender 3
age 104
hypertension 2
heart_disease 2
ever_married 2
work_type 5
Residence_type 2
avg_glucose_level 3979
bmi 418
smoking_status 4
stroke 2
dtype: int64

```
In [8]: stroke_df['gender'].value_counts()# checking gender categories
```

Out[8]: Female 2994
Male 2115
Other 1
Name: gender, dtype: int64

```
In [9]: stroke_df['work_type'].value_counts()
```

Out[9]: Private 2925
Self-employed 819
children 687
Govt_job 657
Never_worked 22
Name: work_type, dtype: int64

In [10]: `stroke_df['smoking_status'].value_counts()`

Out[10]: never smoked 1892
Unknown 1544
formerly smoked 885
smokes 789
Name: smoking_status, dtype: int64

In [11]: `stroke_df['Residence_type'].value_counts()`

Out[11]: Urban 2596
Rural 2514
Name: Residence_type, dtype: int64

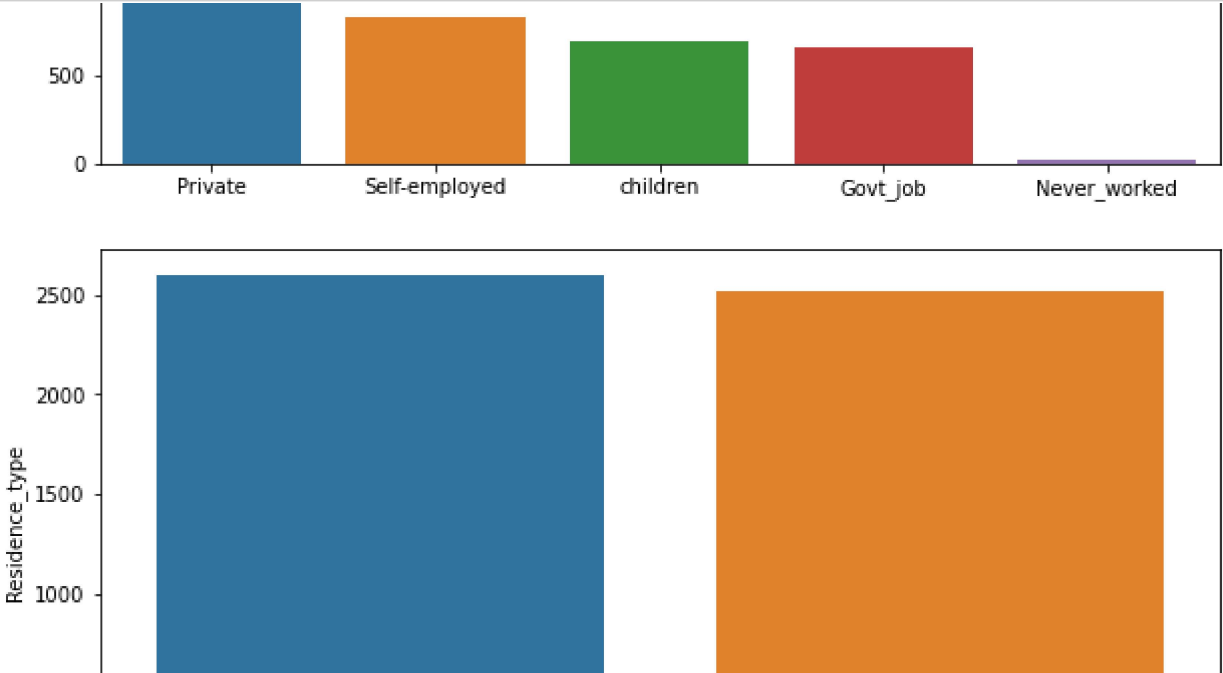
In [12]: `stroke_df.describe(include=['O'])# statistical statistics for categorical variables`

Out[12]:

	gender	ever_married	work_type	Residence_type	smoking_status
count	5110	5110	5110	5110	5110
unique	3	2	5	2	4
top	Female	Yes	Private	Urban	never smoked
freq	2994	3353	2925	2596	1892

Bar plots can be used to compare distributions for categorical variables. We can visualize which subgroups tend to be the most common and how each group compares to one another.

In [13]: `#plotting the categorical variables`
`stroke_barplot = stroke_df[['gender', 'ever_married', 'work_type', 'Residence_type', 'smoking_status','stroke']]`
`for i in stroke_barplot.columns:`
`plt.figure(figsize=(10,5))`
`cat_num = stroke_df[i].value_counts()`
`sns.barplot(x=cat_num.index, y=cat_num)`
`plt.show()`



In [14]: `stroke_df.isna().sum() # check for missing values bmi has 201 missing values 4% of the values are missing`

Out[14]: id 0
gender 0
age 0
hypertension 0
heart_disease 0
ever_married 0
work_type 0
Residence_type 0
avg_glucose_level 0
bmi 201
smoking_status 0
stroke 0
dtype: int64

```
In [15]: ► stroke_df[stroke_df.columns[stroke_df.isna().any()]] # further exploration on the missing bmi values
```

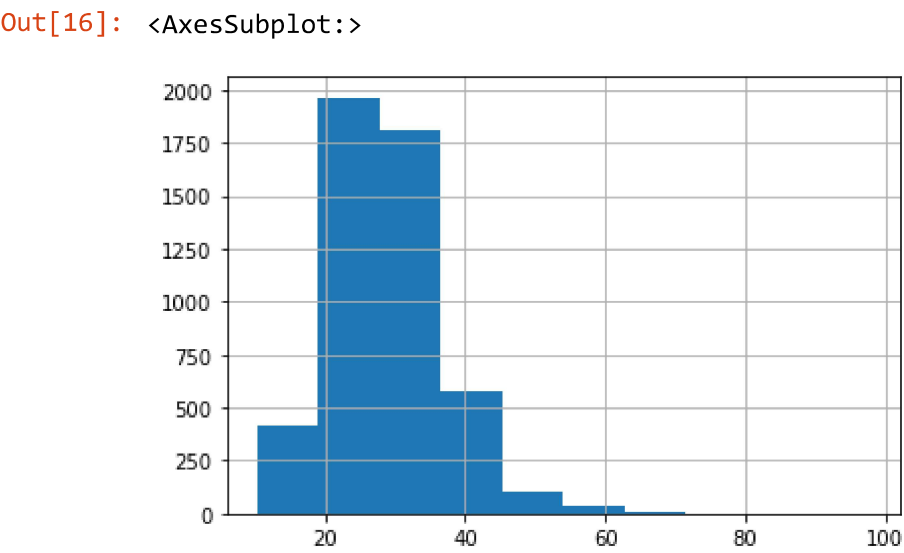
Out[15]:

	bmi
0	36.6
1	NaN
2	32.5
3	34.4
4	24.0
...	...
5105	NaN
5106	40.0
5107	30.6
5108	25.6
5109	26.2

5110 rows × 1 columns

A histogram can be used to visualize the distribution of numerical variables. We use a histogram to look at BMI in order to determine how missing values should be handled. As shown above, BMI has 201 missing values and must be handled appropriately.

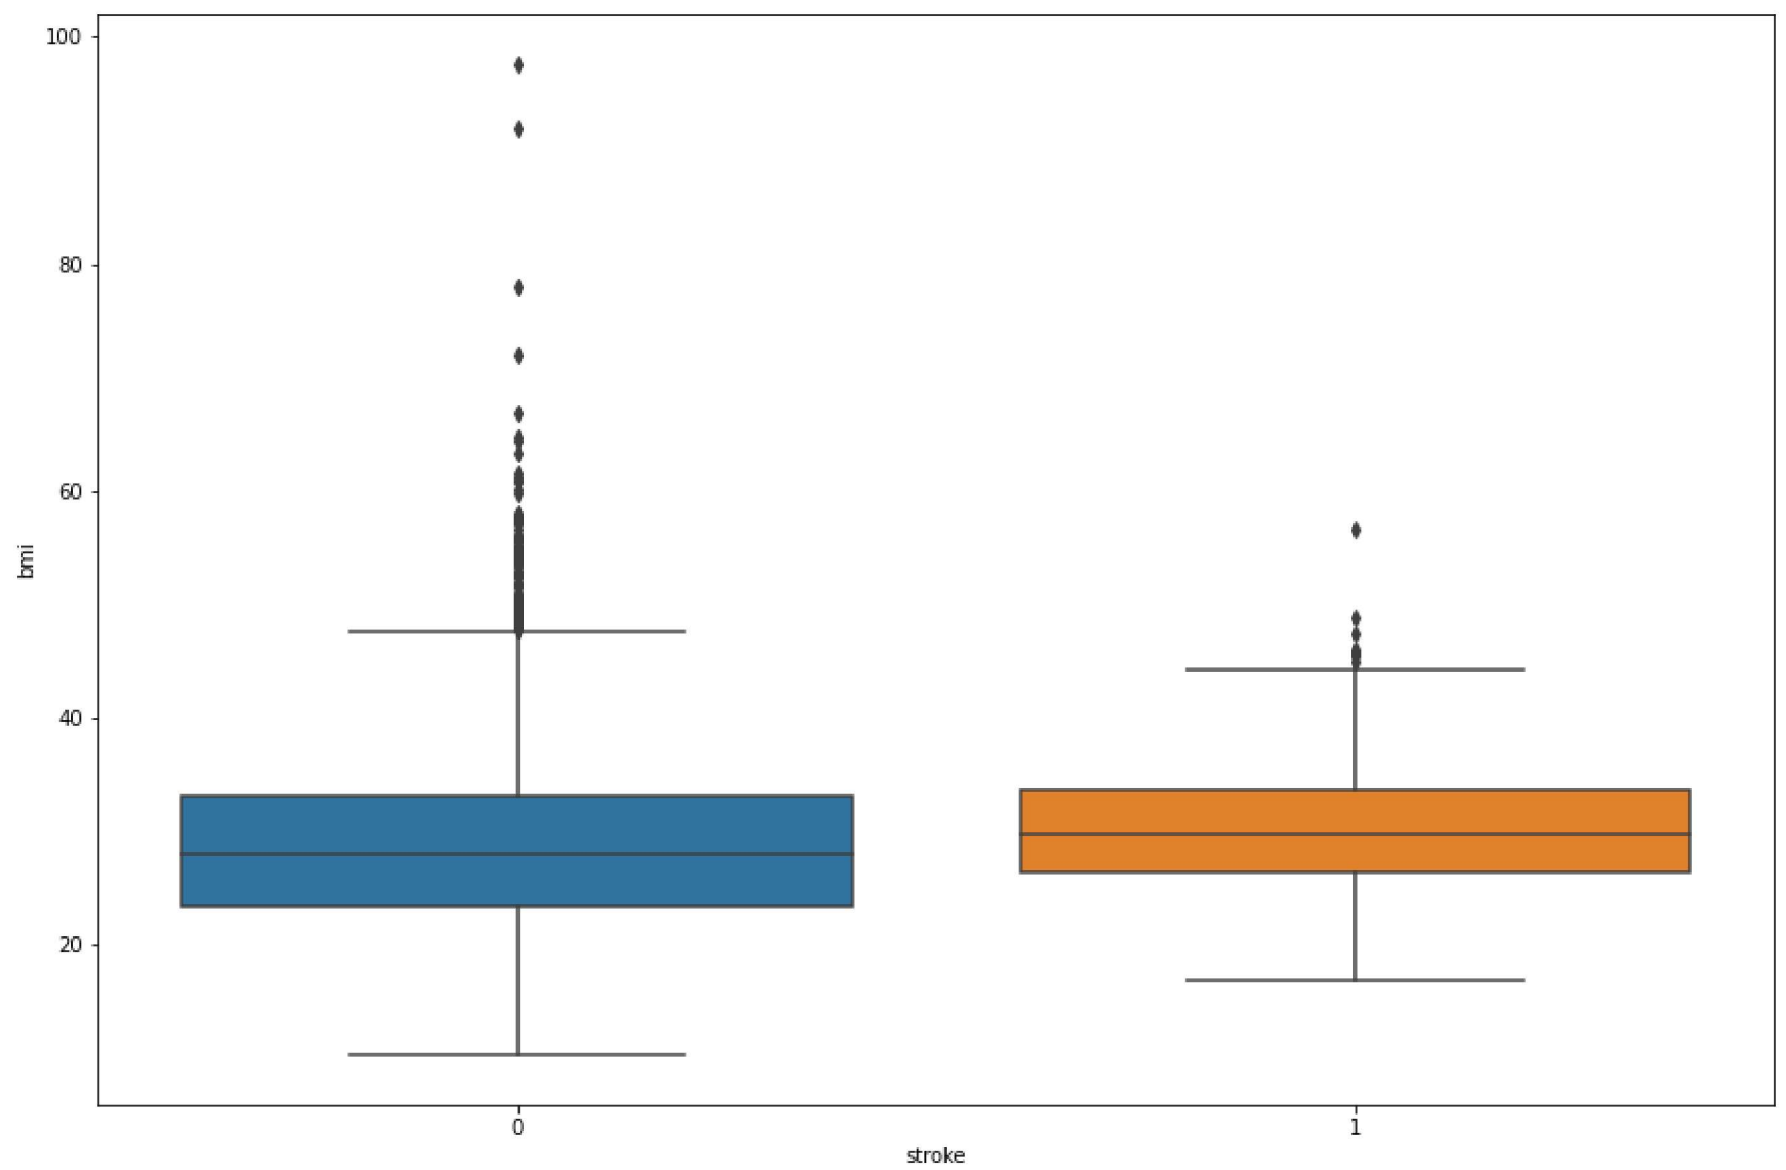
```
In [16]: ► stroke_df['bmi'].hist() #plot histogram for bmi
```



Further analyzing BMI, it appears that while median value for BMI is slightly lower for those who have not suffered from a stroke, there are more outliers and a larger range in values.

```
In [17]: ▶ plt.figure(figsize=(15,10)) # boxplot for bmi
sns.boxplot(x='stroke',y='bmi',data=stroke_df)
```

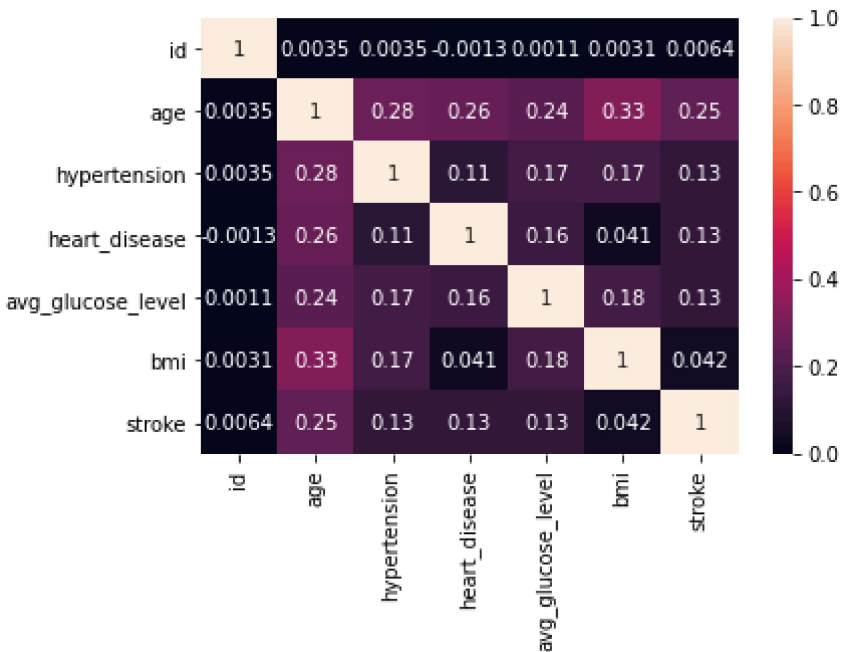
```
Out[17]: <AxesSubplot:xlabel='stroke', ylabel='bmi'>
```



Correlation matrix gives the relationship between our variables and can assist with dimensionality reduction.

```
In [18]: ▶ corrMatrix = stroke_df.corr()
sns.heatmap(corrMatrix, annot=True )
plt.figure(figsize=(45,25))
```

Out[18]: <Figure size 3240x1800 with 0 Axes>



<Figure size 3240x1800 with 0 Axes>

Checking correlation between numerical predictors can assist with dimensionality reduction by removing correlated features

```
In [19]: ▶ stroke_df.corr() #checking correlation between the numerical predictors
```

Out[19]:

	id	age	hypertension	heart_disease	avg_glucose_level	bmi	stroke
id	1.000000	0.003538	0.003550	-0.001296	0.001092	0.003084	0.006388
age	0.003538	1.000000	0.276398	0.263796	0.238171	0.333398	0.245257
hypertension	0.003550	0.276398	1.000000	0.108306	0.174474	0.167811	0.127904
heart_disease	-0.001296	0.263796	0.108306	1.000000	0.161857	0.041357	0.134914
avg_glucose_level	0.001092	0.238171	0.174474	0.161857	1.000000	0.175502	0.131945
bmi	0.003084	0.333398	0.167811	0.041357	0.175502	1.000000	0.042374
stroke	0.006388	0.245257	0.127904	0.134914	0.131945	0.042374	1.000000

```
In [20]: ▶ stroke_df.groupby('stroke').count() # eda on predictor variable
```

Out[20]:

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status
stroke											
0	4861	4861	4861	4861	4861	4861	4861	4861	4861	4700	4861
1	249	249	249	249	249	249	249	249	249	209	249

Following an extensive EDA phase, four features are dropped from the dataset in order to improve model performance and decrease computation time.

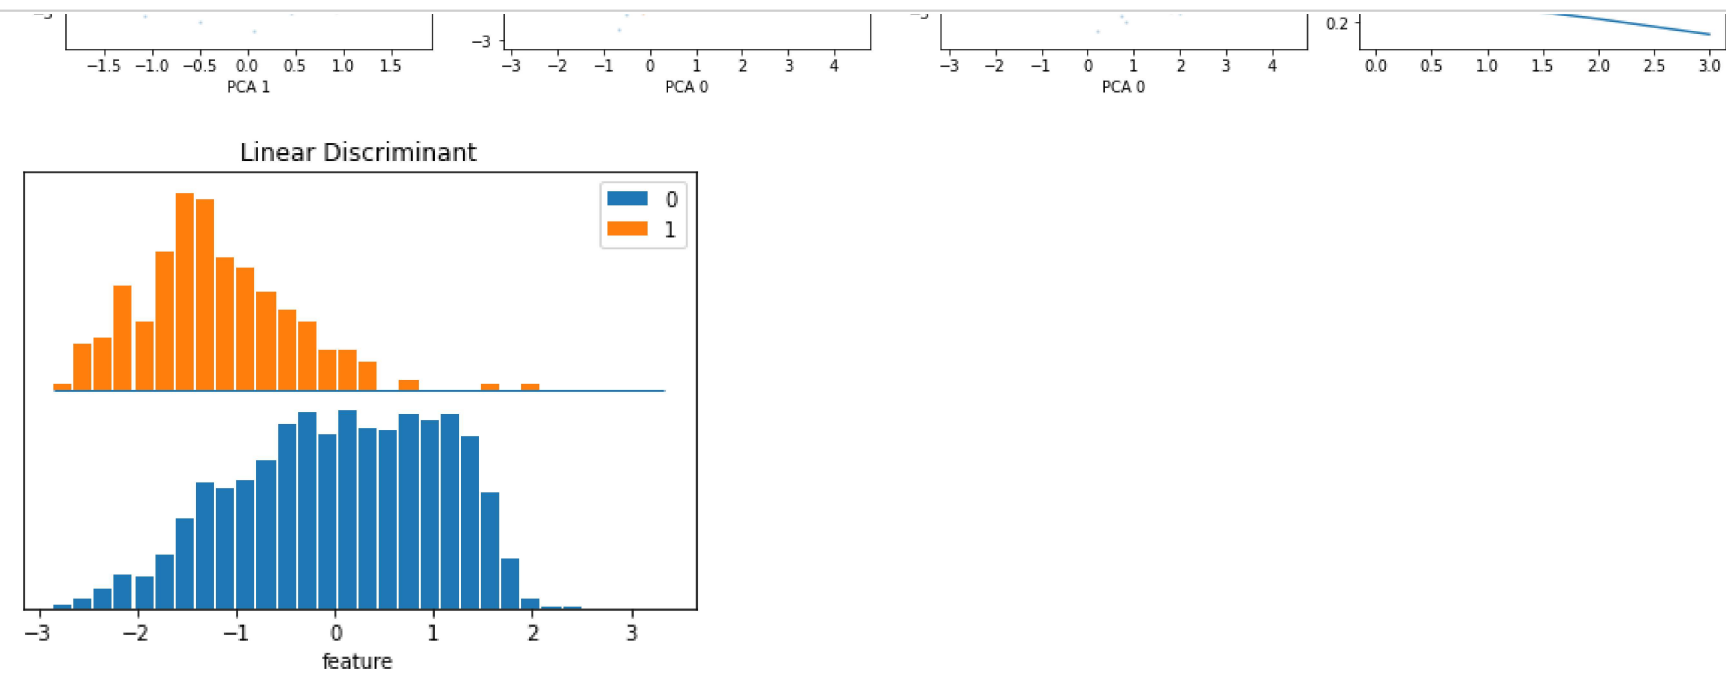
```
In [21]: # further exploration
import dabl
dabl.clean(stroke_df, verbose=2).head(2)
```

Detected feature types:
continuous 4
dirty_float 0
low_card_int 0
categorical 7
date 0
free_string 0
useless 1
dtype: int64

Out[21]:

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status
0	9046	Male	67.0	0	1	Yes	Private	Urban	228.69	36.6	formerly smoked
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	202.21	NaN	never smoked

```
In [22]: # further exploration of the dataset
dabl.plot(stroke_df, 'stroke')
```



```
In [23]: # drop the lifestyle columns from the dataset
stroke_health= stroke_df.drop(columns=['Residence_type', 'work_type', 'ever_married','id'])
stroke_health.head()
```

Out[23]:

	gender	age	hypertension	heart_disease	avg_glucose_level	bmi	smoking_status	stroke
0	Male	67.0	0	1	228.69	36.6	formerly smoked	1
1	Female	61.0	0	0	202.21	NaN	never smoked	1
2	Male	80.0	0	1	105.92	32.5	never smoked	1
3	Female	49.0	0	0	171.23	34.4	smokes	1
4	Female	79.0	1	0	174.12	24.0	never smoked	1

We convert the blood sugar column from numerical to categorical in order to improve signal to noise ratio. As a result, we are able to fit our model according to 4 'bins.' These bins decrease the impact of noise. The four bins are: low, normal, borderline and high.

```
In [24]: #feature engineering Create new column for blood sugar with 4 categories
# stroke_health['Blood_sugar'] = pd.cut(stroke_health.avg_glucose_level, bins=4,
#     labels=["low", "normal", "borderline", "high"]).head()
# stroke_health.head()

stroke_health['Blood_sugar'] = pd.cut(stroke_health.avg_glucose_level, bins=[0,50,72,150,271],
    labels=["low", "normal", "borderline", "high"]).head()
stroke_health.head()
```

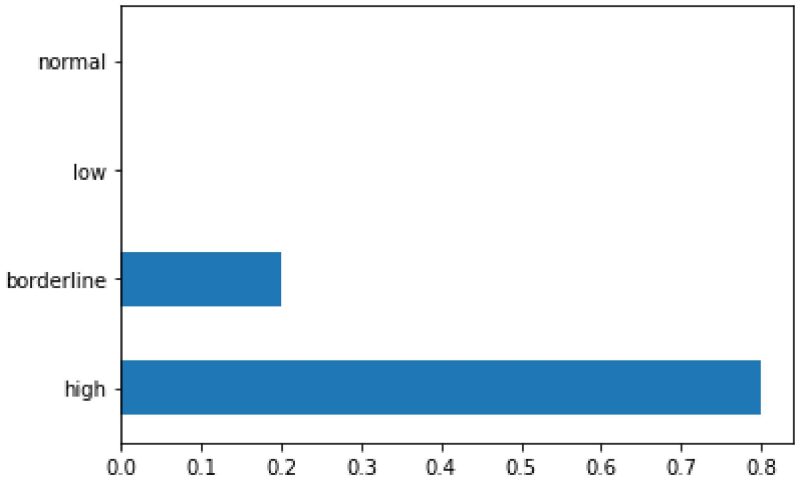
Out[24]:

	gender	age	hypertension	heart_disease	avg_glucose_level	bmi	smoking_status	stroke	Blood_sugar
0	Male	67.0	0	1	228.69	36.6	formerly smoked	1	high
1	Female	61.0	0	0	202.21	NaN	never smoked	1	high
2	Male	80.0	0	1	105.92	32.5	never smoked	1	borderline
3	Female	49.0	0	0	171.23	34.4	smokes	1	high
4	Female	79.0	1	0	174.12	24.0	never smoked	1	high

```
In [25]: ▶ stroke_health['Blood_sugar'].value_counts(normalize=True)
```

```
Out[25]: high      0.8
borderline  0.2
low         0.0
normal      0.0
Name: Blood_sugar, dtype: float64
```

```
In [26]: ▶ #plot blood sugar
stroke_health['Blood_sugar'].value_counts(normalize=True).plot(kind='barh')
plt.show()
```



The dataset is inspected for any missing values which are subsequently imputed by the mode value. This approach is the most representative of the sample and unlike K-Nearest Neighbors, does not yield negative values.

```
In [27]: ▶ # fill missing bmi values with mode
stroke_health['bmi'] = stroke_health['bmi'].fillna(stroke_health['bmi'].mode()[0])
```

```
In [28]: ▶ stroke_health.dtypes
```

```
Out[28]: gender      object
age      float64
hypertension  int64
heart_disease  int64
avg_glucose_level  float64
bmi      float64
smoking_status  object
stroke     int64
Blood_sugar  category
dtype: object
```

Transforming categorical values to numerical, because models can only handle only numeric values. In this instance our categorical values are gender, blood_sugar and smoking_status.

```
In [29]: ▶ le = preprocessing.LabelEncoder()

# Transform the categorical columns for modeling
stroke_health['gender'] = le.fit_transform(stroke_health.gender)
stroke_health['Blood_sugar'] = le.fit_transform(stroke_health.Blood_sugar)
stroke_health['smoking_status'] = le.fit_transform(stroke_health.Blood_sugar)
stroke_health.head()
```

```
Out[29]:
```

	gender	age	hypertension	heart_disease	avg_glucose_level	bmi	smoking_status	stroke	Blood_sugar
0	1	67.0	0	1	228.69	36.6	1	1	1
1	0	61.0	0	0	202.21	28.7	1	1	1
2	1	80.0	0	1	105.92	32.5	0	1	0
3	0	49.0	0	0	171.23	34.4	1	1	1
4	0	79.0	1	0	174.12	24.0	1	1	1

We are now partitioning the dataset into 60% validation and 40% train. Splitting the data will help us determine whether the model's guesses are correct.

```
#partition the dataset into 60% validation and 40% train split

trainData, validData = train_test_split(stroke_health, test_size=0.4, random_state=2) # partition the dataset
print( trainData.shape, validData.shape)

#split the dataset predictor and outcome variables
predictors = list(stroke_health.columns) # predictor variables
outcome = 'stroke'# outcome variable
predictors.remove(outcome)
predictors.remove("avg glucose level") # drop the avg glucose
```



```
print(predictors)
train_X = trainData[predictors]
train_y = trainData[outcome]
valid_X = validData[predictors]
valid_y = validData[outcome]
```

Our first model is Naive Bayes classifier. We chose this model because it is easy and fast to predict the class of the test data set.

```
In [31]: ► # naive bayes model

# run naive Bayes
stroke_nb = MultinomialNB(alpha=0.09)
stroke_nb.fit(train_X, train_y)

# predict probabilities
predProb_train = stroke_nb.predict_proba(train_X)
predProb_valid = stroke_nb.predict_proba(valid_X)

# predict class membership
y_valid_pred = stroke_nb.predict(valid_X)
y_train_pred = stroke_nb.predict(train_X)

# classification summary for the naive Bayes
classificationSummary(train_y, stroke_nb.predict(train_X))

print(classification_report(valid_y, stroke_nb.predict(valid_X), digits=4))
```

Confusion Matrix (Accuracy 0.8604)

		Prediction				
		0	1			
Actual	0	2567	350			
	1	78	71			
		precision	recall	f1-score	support	
0		0.9739	0.8843	0.9269	1944	
1		0.1935	0.5400	0.2850	100	
accuracy				0.8674	2044	
macro avg		0.5837	0.7121	0.6059	2044	
weighted avg		0.9358	0.8674	0.8955	2044	

Naive Bayes model yields 93.58% precision, recall 86.74%, and f-1 score 89.55%.

Our second model is Logistic Regression model. We chose this model because of its simplicity and easy implementation.

```
In [32]: ► #Logistic regression model
#L2 penalty to avoid regularization

logit_reg = LogisticRegression(penalty="l2", C=1e42, solver='liblinear')
logit_reg.fit(train_X, train_y, )

logit_reg_pred_valid = logit_reg.predict(valid_X)
logit_reg_pred_train = logit_reg.predict(train_X)

print('intercept ', logit_reg.intercept_[0])
print(pd.DataFrame({'coeff': logit_reg.coef_[0]}, index=train_X.columns).transpose())
print()
print('AIC', AIC_score(valid_y, logit_reg.predict(valid_X), df = len(train_X.columns) + 1))
```

```
intercept 0.5302911544446849
gender    age  hypertension  heart_disease  bmi \
coeff 0.117552 0.067771 0.501984 0.314333 0.013988

smoking_status  Blood_sugar
coeff -1.98677 -1.98677

AIC -411.3951527038125
```

```
In [33]: ► # classification summary for the logistic regression
classificationSummary(train_y, logit_reg.predict(train_X))

print(classification_report(valid_y, logit_reg.predict(valid_X), digits=4))
```

Confusion Matrix (Accuracy 0.9521)						
Actual		Prediction				
		0	1			
0	2917	0				
1	147	2				
			precision	recall	f1-score	support
0			0.9525	1.0000	0.9757	1944
1			1.0000	0.0300	0.0583	100
accuracy					0.9525	2044
macro avg			0.9762	0.5150	0.5170	2044
weighted avg			0.9548	0.9525	0.9308	2044

Here we are adding class weight to our logistic regression to balance highly imbalanced data.

```
In [34]: ► #Logistic regression model
#add class weights to balance the Y variable to adress class imbalance

bal_reg = LogisticRegression(penalty="l2", C=1e42, solver='liblinear',class_weight='balanced')
bal_reg.fit(train_X, train_y, )

bal_reg_pred_valid = bal_reg.predict(valid_X)
bal_reg_pred_train = bal_reg.predict(train_X)

print('intercept ', bal_reg.intercept_[0])
print(pd.DataFrame({'coeff': bal_reg.coef_[0]}, index=train_X.columns).transpose())
print()
print('AIC', AIC_score(valid_y, bal_reg.predict(valid_X), df = len(train_X.columns) + 1))
```

intercept 6.599042139757904

	gender	age	hypertension	heart_disease	bmi \
coeff	-0.003709	0.06887	0.58086	0.574634	0.014841

	smoking_status	Blood_sugar
coeff	-2.782265	-2.782265

AIC 3168.5451303642963

```
In [35]: ► # classification summary for the balanced logistic regression
classificationSummary(train_y, bal_reg.predict(train_X))

print(classification_report(valid_y, bal_reg.predict(valid_X), digits=4))
```

Confusion Matrix (Accuracy 0.7260)							
		Prediction					
Actual	0	1					
0	2108	809					
1	31	118					
			precision	recall	f1-score	support	
0			0.9860	0.7227	0.8341	1944	
1			0.1292	0.8000	0.2225	100	
accuracy						0.7265	2044
macro avg			0.5576	0.7614	0.5283	2044	
weighted avg			0.9441	0.7265	0.8042	2044	

Balanced Logistic Regression model yields 94.41% precision, recall 72.65%, and f-1 score 80.42%.

Our third model is Random Forest. We chose this model due to its high level of accuracy and ability to perform both regression and classification tasks.

```
In [36]: ▶ #random forest model
rf = RandomForestClassifier(n_estimators=365, random_state=13, class_weight='balanced')
rf.fit(train_X, train_y)
classificationSummary(train_y, rf.predict(train_X))

print(classification_report(valid_y, rf.predict(valid_X), digits=4))
```

Confusion Matrix (Accuracy 0.9958)							
		Prediction					
Actual	0	1					
0	2904	13					
1	0	149					
			precision	recall	f1-score	support	
0			0.9548	0.9887	0.9714	1944	
1			0.2903	0.0900	0.1374	100	
accuracy						0.9447	2044
macro avg			0.6226	0.5393	0.5544	2044	
weighted avg			0.9223	0.9447	0.9306	2044	

Random Forest model yields 92.23% precision, recall 94.47%, and f-1 score 93.06%.

Our last model is Linear Discriminant model. We chose this model due to its ability to reduce dimensionality in data.

```
In [37]: ▶ #Linear discriminant model
lda_reg = LinearDiscriminantAnalysis()
lda_reg.fit(train_X, train_y)
classificationSummary(train_y, lda_reg.predict(train_X))

print(classification_report(valid_y, lda_reg.predict(valid_X), digits=4))
```

Confusion Matrix (Accuracy 0.9491)							
		Prediction					
Actual	0	1					
0	2905	12					
1	144	5					
			precision	recall	f1-score	support	
0			0.9542	0.9959	0.9746	1944	
1			0.4667	0.0700	0.1217	100	
accuracy						0.9506	2044
macro avg			0.7104	0.5329	0.5482	2044	
weighted avg			0.9303	0.9506	0.9329	2044	

Linear Discriminant model yields 93.29% precision, recall 95.06%, and f-1 score 93.29%.

```
In [39]: ▶ print(regressionSummary(valid_y, lda_reg.predict(valid_X)))
```

Regression statistics

Mean Error (ME) : 0.0416

Root Mean Squared Error (RMSE) : 0.2223

Mean Absolute Error (MAE) : 0.0494

None

Based on the abovementioned scores, Linear Discriminant model had highest f-1 score and out performed the other models. Therefore, we recommend using this model for further research and data anylysis.

```
In [ ]: ▶
```